

MASON M. THOMPSON

mason.thompson@berkeley.edu · Berkeley, California · [linkedin.com/in/mmt510](https://www.linkedin.com/in/mmt510) · [ALL-LLM Portfolio](#)

PROFILE

Senior PM with 8+ years building integration, workflow automation, and data infrastructure products across HR, benefits, and SaaS platforms. More recently, I've designed and operated production AI systems spanning speech recognition, speaker diarization, HITL review workflows, and distributed inference architectures. Looking for roles where product ownership and technical depth are equally important.

SKILLS

AI / ML: ASR, Speaker Diarization, Persona Modeling, Human-in-the-Loop (HITL) Pipelines, PyTorch, MLX, Empirical Evaluation, Failure Mode Analysis, Training Data Quality

Systems: Distributed Inference (Mac + PC + GPU hybrid), Streaming Audio (PCM), FastAPI, Tailscale, Structured Logging, GPU Power Monitoring, venv Orchestration

Full Stack: Python, Swift / SwiftUI, React, Node, Supabase, REST APIs, Streamlit, WebSockets

Product: 0→1 Development, API & Systems Design, Roadmapping, Cross-functional Execution, Backlog Management, Operator/Developer Workflows

EXPERIENCE

ALL-LLM POD 06/2025 – Present

Creator — Multimodal AI Chat Platform · Berkeley, CA

- Built the data foundation for a persona-driven AI platform by processing 563+ podcast episodes into structured speaker-labeled training data. Combined WhisperX, NeMo, and custom HITL workflows to detect and remediate transcription, diarization, and attribution failures before they propagated downstream.
- Defined and delivered a multimodal AI platform spanning mobile, speech, inference, and memory systems. Built an iOS client running AI personas over a distributed backend (MLX, GPU, FastAPI, Tailscale) with streaming ASR/TTS, retrieval-augmented memory, and production telemetry.
- Diagnosed a critical silent failure in the ASR decoder corrupting speaker-labeled data at scale; traced the root cause to beam-search and timestamp-token interactions, shipped a hybrid recovery pipeline, and quantified compute-versus-quality tradeoffs.
- Built self-service operator tooling for transcript forensics, quality assurance, and data repair, enabling rapid investigation and remediation of AI-generated errors through auditable playback, review, and correction workflows.

SLOPUME.dev 03/2026 – Present

Creator — AI-Powered Job Search Platform · Berkeley, CA

- Built an AI-native job search platform designed to keep operators in control of the application process. Developed a human-in-the-loop resume rewriting workflow where AI proposes targeted edits and users accept, reject, or refine changes at the individual bullet level.
- Designed and implemented job discovery and prioritization systems aggregating opportunities across Greenhouse, Lever, and Ashby ATS platforms, scoring roles by fit, company quality, and AI-market relevance to compress hours of manual screening into an actionable shortlist.
- Owned end-to-end product development and technical architecture, including FastAPI services, React frontend, Supabase infrastructure, multi-source ingestion pipelines, and AI-assisted analysis workflows — shipping independently from concept through production deployment.

Blissfully 09/2021 – 03/2022

Technical Product Manager, B2B Integrations · New York, NY (remote)

- Owned 0→1 integration strategy for B2B SaaS connectors — defined API schemas, compatibility layers, and developer sandbox requirements before the engineering team was formed.
- Authored PRDs and technical specs establishing system boundaries, data contracts, and execution roadmap from first principles.
- Led nearshore team of 3 engineers to ship multiple integrations in 8 weeks; contributed to pre-acquisition product readiness.

Ease 03/2020 – 04/2021

Product Manager, Integrations · San Francisco, CA

- Owned end-to-end health insurance integrations (0→1 → scale) for an HRIS platform serving 70K+ SMBs and 2M+ end users — built core infrastructure for carrier connectivity.

MASON M. THOMPSON

mason.thompson@berkeley.edu · Berkeley, California · [linkedin.com/in/mmt510](https://www.linkedin.com/in/mmt510) · [ALL-LLM Portfolio](#)

- Led 8-engineer team; automated audit workflows for a 10-person ops team, cutting weekly task volume 93% and eliminating a 3-month backlog.
- Launched 5 carrier integrations in 90 days, adding 300K+ end users and ~\$2M projected ARR.

Zenefits 09/2015 – 02/2019

Sr. Product Operations Specialist (promoted 2x: Analyst → Specialist → Sr. Specialist) · San Francisco, CA

- Designed and launched the Universal Enrollment File (UEF) — a unified data layer enabling interoperability across carrier formats, powering 43% of enrollment workflows.
- Automated 74% of Open Enrollment processing within 5 months; reduced enrollment error rates from 8% to <1% via carrier integrations and multi-tenant portal improvements.
- Built and scaled a Google Sheets-based bug tracking and triage system during a major React migration, cutting reported bugs 80% and improving visibility for 100+ offshore ops staff.
- Led real-time iframe enrollment integration (Unum) enabling instant decisioning without direct handling of sensitive PHI.

EDUCATION & CERTIFICATIONS

University of California, Berkeley — B.A. Molecular & Cell Biology, Concentration in Biochemistry & Molecular Biology

NVIDIA — Generative AI & LLM Associate Certification, February 2026